

CS5473 - Project 5

Author: Enzo Durel

Professor: Dr. Chongle Pan

April 11, 2025



Contents

| | | |
|----------|-----------------------|----------|
| 1 | Problem 1 | 1 |
| 1.1 | Runtimes | 1 |
| 1.2 | Same Node | 1 |
| 1.3 | Diff Node | 1 |
| 1.4 | Linear Regression | 1 |
| 2 | Problem 2 | 3 |
| 2.1 | Wall-clock Time Table | 3 |
| 2.2 | Speedup | 3 |
| 2.3 | Efficiency | 3 |
| 2.4 | Discussion | 3 |
| 3 | Problem 3 | 4 |
| 3.1 | Runtimes | 4 |
| 3.2 | Speedup | 4 |
| 3.3 | Efficiency | 4 |
| 3.4 | Discussion | 4 |
| 4 | Problem 4 | 5 |
| 4.1 | Runtimes | 5 |
| 4.2 | Speedup | 5 |
| 4.3 | Efficiency | 5 |
| 4.4 | Discussion | 5 |

List of Figures

| | | |
|---|--|---|
| 1 | Linear Regression for same | 1 |
| 2 | Linear Regression for diff | 2 |
| 3 | Linear Regression for both methods | 2 |

List of Tables

| | | |
|----|---|---|
| 1 | Problem 1 Runtime | 1 |
| 2 | Problem 2 Wall-clock Time Table | 3 |
| 3 | Problem 2 Speedup Table | 3 |
| 4 | Problem 2 Efficiency Table | 3 |
| 5 | Problem 3 Runtimes | 4 |
| 6 | Problem 3 Speedup | 4 |
| 7 | Problem 3 Efficiency | 4 |
| 8 | Problem 4 Runtimes | 5 |
| 9 | Problem 4 Speedup | 5 |
| 10 | Problem 4 Efficiency | 5 |

1 Problem 1

1.1 Runtimes

Table 1: Problem 1 Runtime

| Integers | 1M | 2M | 4M | 8M |
|----------|--------------|--------------|--------------|--------------|
| Same RT | 0.0002394965 | 0.0004862981 | 0.0012506239 | 0.0032828204 |
| Diff RT | 0.0004300362 | 0.0008670914 | 0.0017082404 | 0.0039351337 |

1.2 Same Node

- Latency $\approx -3.488867e - 04s$
- Bandwidth $\approx 9016.07MB/s$

1.3 Diff Node

- Latency $\approx -1.498357e - 04s$
- Bandwidth $\approx 7957.72MB/s$

1.4 Linear Regression

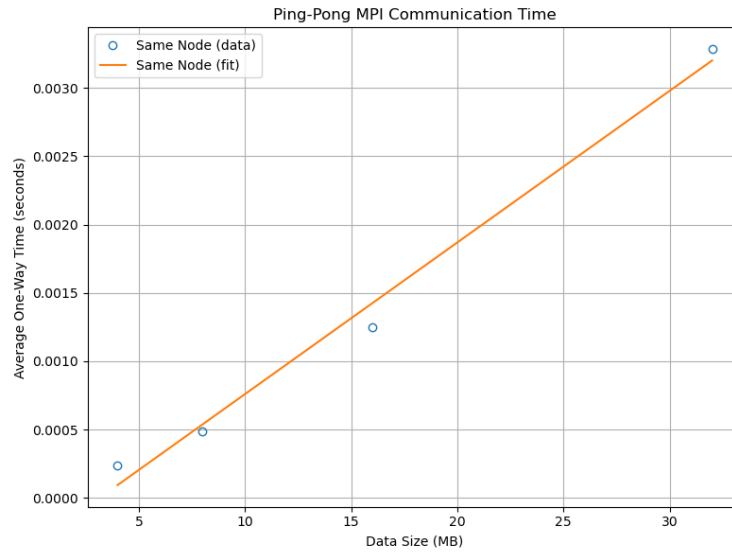


Fig. 1: Linear Regression for same




Fig. 2: Linear Regression for diff

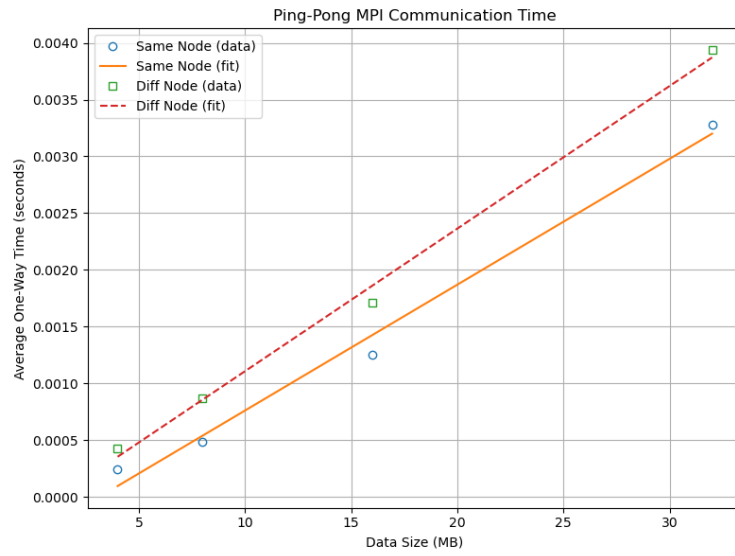


Fig. 3: Linear Regression for both methods

2 Problem 2

2.1 Wall-clock Time Table

Table 2: Problem 2 Wall-clock Time Table

| Array size | 262144 | 524288 | 1048576 |
|-------------|-------------|-------------|-------------|
| serial | 0.000902495 | 0.001672294 | 0.003332579 |
| 2 processes | 0.002470 | 0.005453 | 0.008972 |
| 4 processes | 0.003374 | 0.004037 | 0.008923 |
| 8 processes | 0.003583 | 0.003975 | 0.008298 |

2.2 Speedup

Table 3: Problem 2 Speedup Table

| Array size | 262144 | 524288 | 1048576 |
|-------------|--------|--------|---------|
| 2 processes | 0.3652 | 0.3067 | 0.3714 |
| 4 processes | 0.2673 | 0.4143 | 0.3733 |
| 8 processes | 0.2518 | 0.4205 | 0.4016 |

2.3 Efficiency

Table 4: Problem 2 Efficiency Table

| Array size | 262144 | 524288 | 1048576 |
|-------------|--------|--------|---------|
| 2 processes | 0.1826 | 0.1534 | 0.1857 |
| 4 processes | 0.0668 | 0.1036 | 0.0933 |
| 8 processes | 0.0315 | 0.0526 | 0.0502 |

2.4 Discussion

Speedup is sub-linear due to communication overhead in MPI. Efficiency drops with increasing process count. With 8 processes and 262144 elements, we've got around 3% efficiency, which suggests that parallel overhead dominates the computation.

The program shows poor scalability especially for small vector sizes. An MPI limitation here is, for lightweight computations, the communication overhead quickly outweighs parallel benefits. Scalability improves slightly with larger input sizes but remains insufficient.

3 Problem 3

3.1 Runtimes

Table 5: Problem 3 Runtimes

| | | | |
|----------------------------------|----------|----------|----------|
| Array size | 262144 | 524288 | 1048576 |
| 4 processes on the same node | 0.019300 | 0.040307 | 0.084053 |
| 4 processes on 4 different nodes | 0.026653 | 0.042971 | 0.112452 |

3.2 Speedup

Table 6: Problem 3 Speedup

| | | | |
|----------------------|--------|--------|---------|
| Array size | 262144 | 524288 | 1048576 |
| Diff vs Same Speedup | 0.7245 | 0.9370 | 0.7474 |

3.3 Efficiency

Table 7: Problem 3 Efficiency

| | | | |
|----------------------------|--------|--------|---------|
| Array size | 262144 | 524288 | 1048576 |
| Efficiency (Diff Node, /4) | 0.1811 | 0.2343 | 0.1868 |

3.4 Discussion

MPI Merge Sort shows better scalability on the same node due to reduced communication latency and faster memory access.

Performance drops on different nodes, see in timing and efficiency metrics. This aligns with the other problems results, meaning that distributed-node latency and bandwidth penalty strongly algorithms. Moreover, merge sort algorithm has a significant data exchange.

4 Problem 4

4.1 Runtimes

Table 8: Problem 4 Runtimes

| | | | | | |
|------------|----------|----------|----------|----------|----------|
| Array size | 1 | 2 | 4 | 8 | 16 |
| Runtime | 0.002640 | 0.001349 | 0.000798 | 0.000388 | 0.000290 |

4.2 Speedup

Table 9: Problem 4 Speedup

| | | | | | |
|-----------|--------|--------|--------|--------|--------|
| Processes | 1 | 2 | 4 | 8 | 16 |
| Speedup | 1.0000 | 1.9570 | 3.3075 | 6.8041 | 9.1034 |

4.3 Efficiency

Table 10: Problem 4 Efficiency

| | | | | | |
|------------|-------|-------|-------|-------|-------|
| Processes | 1 | 2 | 4 | 8 | 16 |
| Efficiency | 1.000 | 0.978 | 0.827 | 0.850 | 0.569 |

4.4 Discussion

The Monte Carlo pi estimation program is strongly scalable. There is small communication overhead, practical speedup around the theoretical speedup and a high efficiency.